

10. NoSQL

- Introducció
- Bases de dades SQL
- Bases de dades NoSQL
- Bases de dades NewSQL



NoSQL: El nom

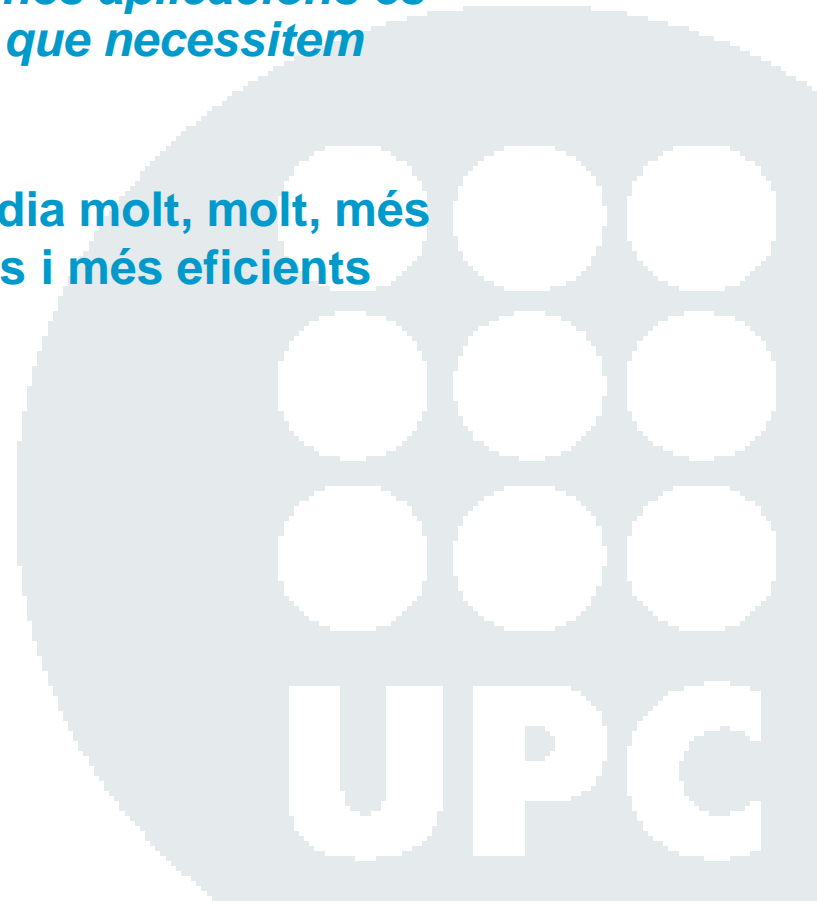
- “SQL” = SGBD relacionals tradicionals
- S’acaba la tendència “One size fits all”:
No tots els problemes de gestió/anàlisi de dades es resolen millor usant **exclusivament** SGBD relacionals tradicionals
- “No SQL” \neq No usar SQL
- “NoSQL” = “Not only SQL” = No usar **només** SGBD tradicionals

No tots els problemes de gestió/anàlisi de dades es resolen millor usant SGBD relacionals tradicionals

- Persistent
- Eficient
- Massiu
- Multi-usuari
- Segur
- Fiable
- Convenient

Per algunes aplicacions és més del que necessitem

Avui en dia molt, molt, més massives i més eficients



SQL Characteristics

- Data stored in columns and tables
- Relationships represented by data
- Data Manipulation Language
- Data Definition Language
- Transactions
- Abstraction from physical layer

SQL Physical Layer Abstraction

- Applications specify what, not how
- Query optimization engine
- Physical layer can change without modifying applications
 - Create indexes to support queries
 - In Memory databases



Data Definition Language

- Schema defined at the start
- Create Table (Column1 Datatype1, Column2 Datatype 2, ...)
- Constraints to define and enforce relationships
 - Primary Key
 - Foreign Key
 - Etc.
- Triggers to respond to Insert, Update , & Delete
- Stored Modules
- Alter ...
- Drop ...
- Security and Access Control

Data Manipulation Language (DML)

- Data manipulated with Select, Insert, Update, & Delete statements
 - Select T1.Column1, T2.Column2 ...
From Table1, Table2 ...
Where T1.Column1 = T2.Column1 ...
- Data Aggregation
- Compound statements
- Functions and Procedures
- Explicit transaction control

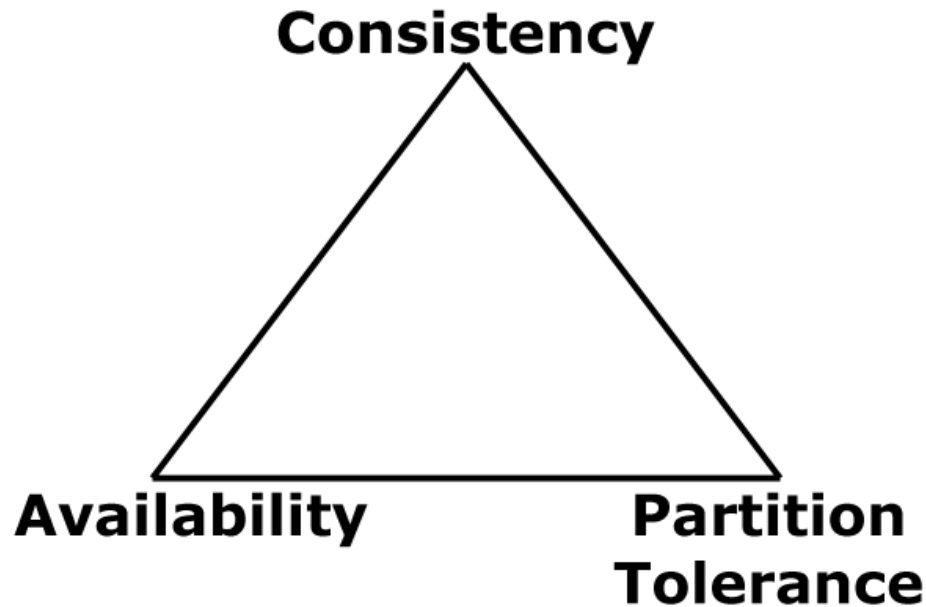
Transactions – ACID Properties

- **Atomic** – All of the work in a transaction completes (commit) or none of it completes
- **Consistent** – A transaction transforms the database from one consistent state to another consistent state. Consistency is defined in terms of constraints.
- **Isolated** – The results of any changes made during a transaction are not visible until the transaction has committed.
- **Durable** – The results of a committed transaction survive failures

NoSQL Distinguishing Characteristics

- Large data volumes
 - Google's "big data"
- Scalable replication and distribution
 - Potentially thousands of machines
 - Potentially distributed around the world
- Queries need to return answers quickly
- Mostly query, few updates
- Asynchronous Inserts & Updates
- Schema-less
- ACID transaction properties are not needed – BASE
- CAP Theorem
- Open source development

CAP THEOREM



The [CAP Theorem](#) by Eric Brewer.

CAP stands for **C**onsistency, **A**vailability, (network) **P**artition tolerance and the theorem claims that *in a distributed system, when there is an inevitable network partition (and the cluster breaks into two or more “islands”), you can’t guarantee both availability (for updates) and consistency.*

BASE Transactions

- Acronym contrived to be the opposite of ACID
 - **B**asically **A**vailable,
 - **S**oft state,
 - **E**ventually **C**onsistent
- Characteristics
 - Weak consistency – stale data OK
 - Availability first
 - Best effort
 - Approximate answers OK
 - Aggressive (optimistic)
 - Simpler and faster

A (Tentative) Classification of NOSQL Databases

■ Key-Value Stores

- Origin: Massive web-based distributed systems (e.g., Amazon, Google, Facebook, etc.)
- Goal: Scalability, efficiency, fault-tolerance, support for unstructured data
- Data Model: Collection of <Key, Value> pairs
- Example: BigTable (HBase), Cassandra, Voldemort, SimpleDB / Dynamo, etc.

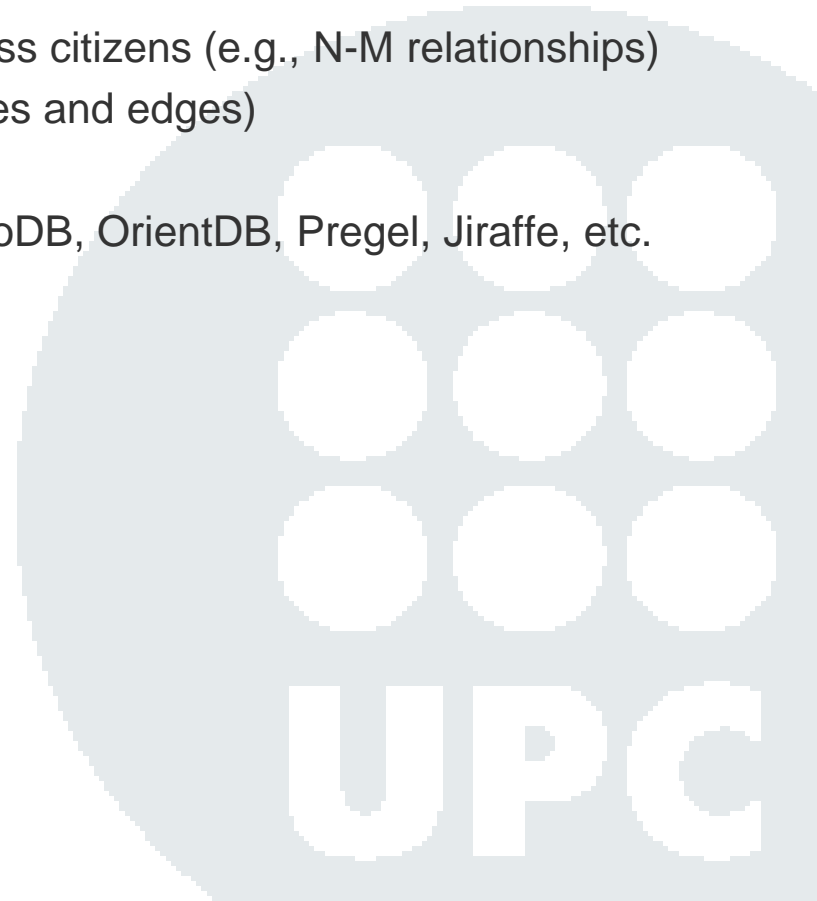
■ Document-stores

- Origin: An evolution of key-value stores where the value is semi-structured
- Goal: Scalability, efficiency, fault-tolerance, support for semi-structured data
- Data Model: Collection of <Key, Doc> pairs, where Doc is typically a XML or JSON document
- Example: MongoDB (JSON), CouchDB (JSON), eXisT (XML), etc.

A (Tentative) Classification of NOSQL Databases

■ Graph Databases

- Origin: Occurrence-based systems, where most queries refer to individuals (e.g., *traverse* queries)
- Goal: Deal with relationships as first-class citizens (e.g., N-M relationships)
- Data Model: Graph structures (i.e., nodes and edges)
- Architecture: Diverse
- Example: Neo4J, DEX, Virtuoso, ArangoDB, OrientDB, Pregel, Jiraffe, etc.



Summary

- SQL Databases
 - Predefined Schema
 - Standard definition and interface language
 - Tight consistency
 - Well defined semantics
- NoSQL Database
 - No predefined Schema
 - Per-product definition and interface language
 - Getting an answer quickly is more important than getting a correct answer

NewSQL

Give Up SQL?

- Compiler translates SQL at compile time into a sequence of low level operations
- Similar to what the NoSQL products make you program in your application
- 30 years of RDBMS experience
 - + Hard to beat the compiler
 - + High level languages are good (data independence, less code, ...)
 - + **Stored** procedures are good!
 - One round trip from app to DBMS rather than one one round trip per record
 - Move the code to the data, not the other way around

NewSQL

Give Up ACID

- If you need data consistency, giving up ACID is a decision to tear your hair out by doing database “heavy lifting” in user code
- Can you guarantee you won’t need ACID tomorrow?



ACID = goodness, in spite of what these guys say

■ Column-Oriented Databases

- Origin: Query-oriented systems (e.g., data warehouses)
- Goal: Improve the valid read ratio
- Data Model: Relational
- Example: HP Vertica, MonetDB, C-Store, etc.

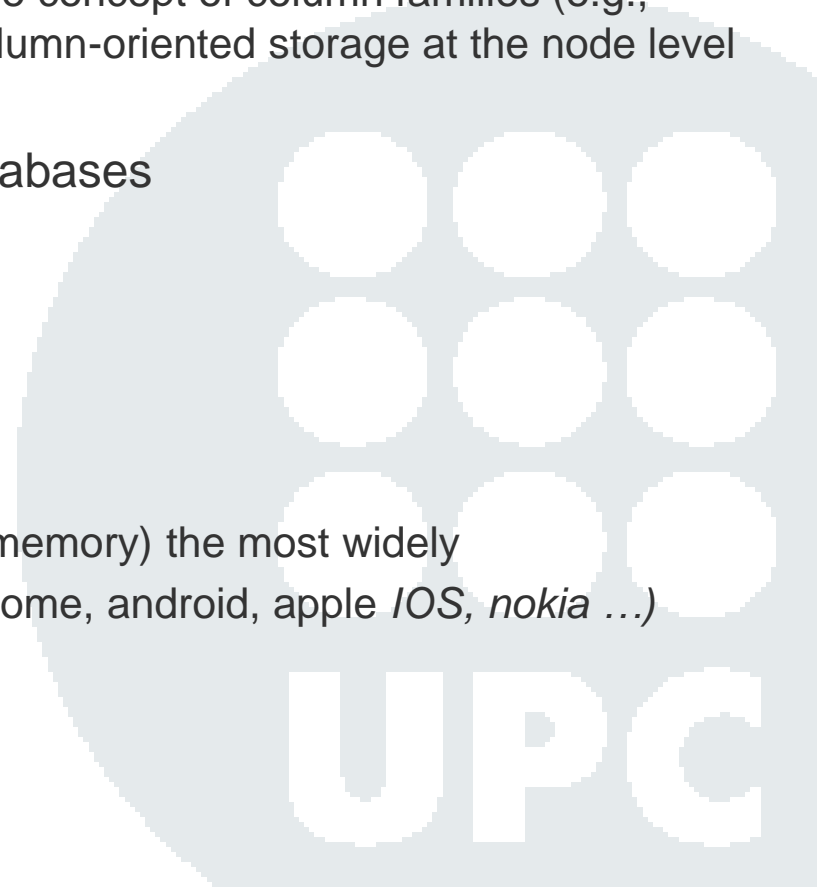
* Note: Some key-value stores introduce the concept of column families (e.g., HBase, Cassandra), which implies local column-oriented storage at the node level

■ In-memory Databases / Embedded Databases

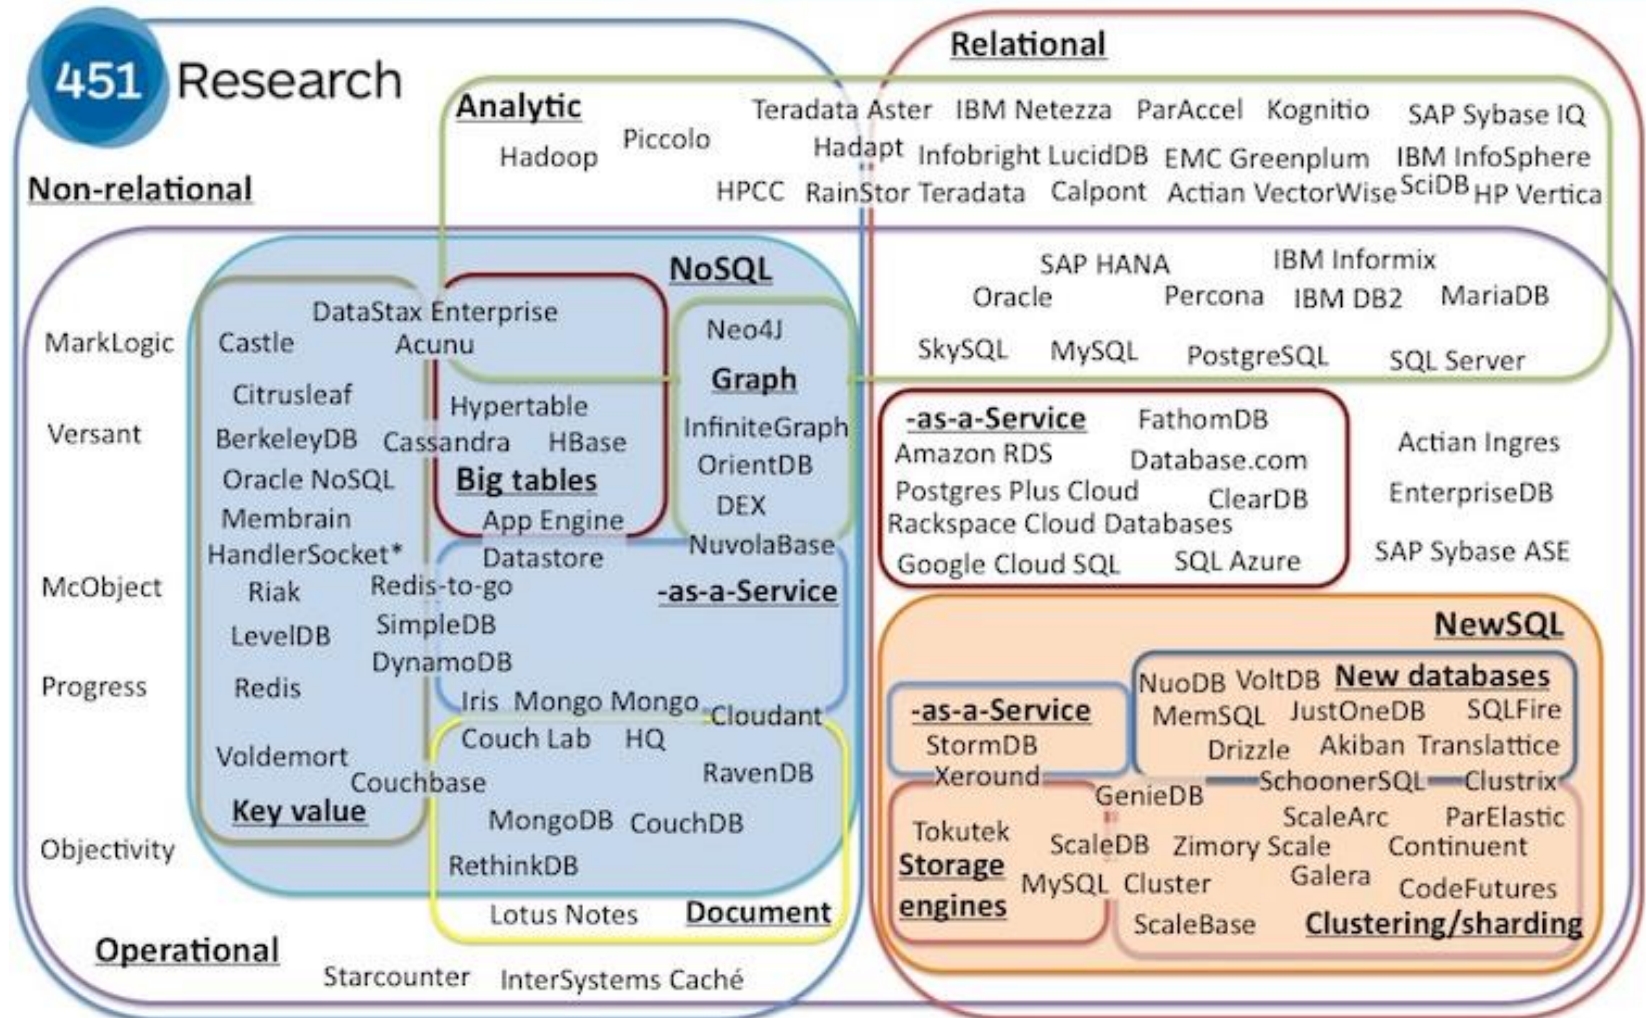
- Origin: SQL is not the problem
- Goal: Improve performance
- Data Model: Relational
- Example:

VoltDB (in-memory),

SQLite (embedded, optionally also in-memory) the most widely deployed database engine (firefox, chrome, android, apple *IOS*, *nokia* ...)



The evolving database landscape



HOW TO WRITE A CV



Leverage the NoSQL boom

- Keith Hare, A comparison of SQL and NoSQL database <http://www.slideshare.net/Muratakal/rdbms-vs-nosql-15797058>

- The 451 Group. The database landscape.
http://blogs.the451group.com/information_management/2012/11/02/updated-database-landscape-graphic/

Michael Stonebraker, [Samuel Madden](#), [Daniel J. Abadi](#), [Stavros Harizopoulos](#), [Nabil Hachem](#), [Pat Helland](#): **The End of an Architectural Era (It's Time for a Complete Rewrite)**. [VLDB 2007](#): 1150-1160

- Michael Stonebraker: **Stonebraker on NoSQL and enterprises**. [Commun. ACM 54](#)(8): 10-11 (2011)

- Michael Stonebraker: **New opportunities for New SQL**. [Commun. ACM 55](#)(11): 10-11 (2012)